FOR SERIES A FOUNDERS AND THEIR CTOS

The Series A AI Growth Kit



Choose the AI project worth your next engineering investment.

Compare six AI opportunities, choose one to investigate, and define a focused proof of value.

Bring four inputs

Start with your growth priority, up to three ideas, a recent baseline, and a workflow owner. Assign someone to resolve any missing input.

01 Explore

Review six opportunities on pages 2 and 3.

02 Compare

Score your top three ideas on page 4.

03 Model

Use the value worksheet on page 5.

04 Scope

Design the test and complete your brief on pages 6 and 7.

Use these planning tools independently or with an implementation partner. Numerical examples are illustrative, not promised returns.

01 / FIND THE OPPORTUNITY

Where AI could support growth

Start with a constraint your team is already trying to remove.

01 Help buyers complete an evaluation

A sales copilot drafts source-linked answers from approved product and security information for review.

Measure response time and accuracy, then conversion for comparable opportunities.

Keep customer commitments and pricing subject to approval.

02 Help new users reach first value

An onboarding agent guides users through one setup task and routes blockers to a person.

Measure completion and time to first value before assessing trial-to-paid conversion.

Require approval for account changes.

03 Resolve a recurring support request

A support agent resolves recurring requests using approved information and permitted actions.

Track verified resolution, escalations, and cost, including failed attempts and human handling.

Assess retention over a longer period.

01 / FIND THE OPPORTUNITY

Where AI could create more capacity

04 Add a paid workflow to your product

An embedded AI feature helps customers complete an analysis or report using their authorized data.

Measure task success, repeat use, willingness to pay, and cost per completed task.

Confirm paid demand before projecting expansion revenue.

05 Increase customer-operations throughput

An operations agent checks cases across permitted systems, prepares actions, and routes exceptions for review.

Measure throughput, handling time, rework, and cost per accepted case.

Distinguish usable capacity from cash savings.

06 Reduce repeatable engineering work

AI-assisted review and quality assurance help engineers check work against relevant specifications.

Measure review time, defects, and engineering effort to understand delivery capacity.

Keep approval and deployment authority with your engineers.

02 / PRIORITIZE

Compare your top three ideas

Score each criterion from 0 to 2. Use 0 for missing or unknown evidence, 1 when more work is needed, and 2 when the condition below is met.

Criterion	Award 2 points when...
Business consequence	The idea supports a current growth or margin priority.
Frequency or reach	The work repeats at a known, meaningful volume.
Measurability	A baseline, owner, and comparison method exist.
Data and integration access	Representative inputs and approved access are ready.
Ownership and scope	Business and technical owners agree on a narrow task.

Candidate	Five scores	Total / 10	Evidence or main gap
A			
B			
C			

Illustrative example: Trial onboarding scores 9 out of 10, but billing access still needs approval. Resolve that dependency before testing.

> Scores prioritize discovery; they do not predict returns. Confirm data permission, a safe test, human review, and a decision owner. Critical security or legal issues override any score.

03 / MAKE THE BUSINESS CASE

Put value and cost on one page

Revenue contribution

Multiply eligible opportunities by the absolute conversion improvement and gross profit per added customer. Subtract AI operating costs for the same period.

Capacity or cash

Multiply adopted tasks by human minutes saved, including review and rework, then divide by 60. Do not count the same hours as both capacity and cash savings.

Trial onboarding, an illustrative example

Assume 500 trials each month, 20% conversion, and $100 gross profit from each added customer's first billed month. AI operations cost $1,800 monthly, including failures, review, monitoring, and support.

Monthly scenario	Low	Base	High
Assumed conversion	23%	25%	28%
Additional paid customers	15	25	40
Additional first-month gross profit	$1,500	$2,500	$4,000
AI operating costs	$1,800	$1,800	$1,800
Net monthly contribution	**-$300**	**$700**	**$2,200**

This first-month-only model excludes later recurring margin, ramp-up, churn, taxes, and discounting. These are assumptions, not measured results or Avahi pricing.

Your inputs	Value and evidence source
Eligible volume and baseline	
Expected improvement	
Benefit and recurring costs	
One-time build, integration, and internal effort	

04 / DESIGN THE TEST

Define what the AI must prove

Choose the simplest approach that meets the need, including non-AI automation. [2]

Approach	Set the boundary
AI assistant	A person approves and carries out the action.
Fixed AI workflow	Code controls predefined steps and permitted actions.
Bounded agent	Limit tools, data, permissions, steps, spending, and escalation.

Set four acceptance gates

Outcome

Compare one primary metric with its documented baseline.

Quality

Define a successful task and count review and rework.

Economics

Include every attempted run and human handling in the cost.

Control

Test permissions, escalation, logs, and failure recovery.

Agree the sample, comparison, and observation window before testing. Keep evaluation cases separate from tuning.

Stop for unauthorized actions, cross-account data exposure, or unresolved critical risks.

Allow enough time for the outcome

Revenue and retention require enough completed cohorts. Until that evidence exists, treat projected gains as hypotheses.

05 / TAKE THIS TO YOUR CTO

Your proof of value brief

Complete one brief. Assign an owner and next action to every unknown.

Decision field	Your working brief
Workflow	What task, input, and accepted output will you test?
Baseline	What is the current metric, period, and source?
Target	What improvement would justify further investment?
Scope	What is excluded, and which simpler option will you compare?
Data and permissions	What may the AI access or change? Who approves?
Evaluation	What sample, comparison, and duration will establish the result?
Cost and control	What investment, operating limit, and acceptance gates apply?
Owners	Who owns the outcome, technical work, and decision?

Go when agreed gates are met. **Modify** the plan to address a fixable gap. **Stop** when costs or unresolved risks outweigh value. Insufficient evidence means more observation.

06 / PUT IT TO WORK

Build your proof of value

Avahi builds and implements AI and agentic AI solutions.

01 Agree one workflow, baseline, and value hypothesis.

02 Build a working solution in your AWS environment.

03 Measure business impact, quality, human effort, cost, and risk.

04 Make a go, modify, or stop recommendation.

Most engagements are planned for **4 to 6 weeks**, depending on scope and readiness. Revenue and retention may need longer observation.

A documented example: myRiva

myRiva reported about a **40 to 50% gain in development efficiency** after Avahi established Amazon Bedrock and connected QA to Confluence. This customer result is not a forecast. [1]

Bring your brief to a Value Workshop

Discuss scope, data readiness, and success criteria using your completed brief.

Request a Value Workshop

AWS funding may be available for qualified projects, subject to approval. Coverage and any customer-paid portion are confirmed before work begins.

AWS Premier Tier Services Partner [3]

[1] myRiva customer story [2] Building effective agents [3] Avahi partnerships

Examples are illustrative unless attributed to a customer.